

CAP theorem >>

Live coding of kafka streams >>> custom serde >><
<https://livebook.manning.com/book/kafka-streams-in-action/chapter-3/165>

<https://www.ibm.com/cloud/learn/cap-theorem>

<https://www.youtube.com/watch?v=7Faly8jORlw>

Very good ppt

https://www.youtube.com/redirect?event=video_description&redir_token=QUFFLUhqa3BxTUtoYmkyLWRnS0c1djRRM3dCSVNnVWEyd3xBQ3Jtc0ttVGhmdWxCV05od2RBdGsxcmNMTHZlWFNpOTNlcmNEekI4N05tSEV4V2ZlV2IfY1FCc05mNk1wOnA0YTFaMXdNRHRLb0JPUiFrMVZYRFY4TFdNeDVoNlIPZlIlalBySTdvWnUwOUhra1Z3ek1meFBmOA&q=https%3A%2F%2Fwww.slideshare.net%2Fpivotal%2Frabbitmq-kafka

kafka joint streaming

when you join the stream make sure you have same key and partition count is also same

rabbit mq

direct exchange > based on exact matching of routing key
topic exchange > wild card of routing key matches for eg ". eu or eu.de.*
fanout exchange: which delivers message to all queues regardless of routing key or pattern matching

tips for kafka

use ack=0

ack =1 to balance between reliability and latency

linger of 5ms is a good thumb rule

increase batch size for higher throughput

low linger for low latency

if you have a large files

> put files in a shared location and send location of files on kafka
break down file to right size and use keys to ensure ordering processing

using keys>

using null keys gives best performance and balanced partition across cluster
to establish ordering you cannot have null keys as ordering will be accomplished based on keys

use keys only if you need ordered messages in real time or joins across different topics

extremely large messages may block consumers if the consumer is not adequate buffer size. if there are some large messages, consider using separate topic.

watch out for zombie brokers in older kafka table

restarting cluster with large number of partitions the leader election takes time.

if consumer crashes or unable to send heartbeats, the partition reassignment will take time and during this time no consumer in the group can process the message